UniAPL: A Unified Adversarial Preference Learning Framework for Instruct-Following

Faqiang Qian[1], **Weikun Zhang**[1], **Ziliang Wang**[1], **Kang An**[1], **Xuhui Zheng**[1],
Liangjian Wen[2] , **Mengya Gao**[1], **Yong Dai**[3], **Yichao Wu**[1*]
[1]Imvision Innovation
[2]Southwest University of Finance and Economics
[3]X-Humanoid
{jingleqian, wayconzhang,elsagao, cyberwu}@ai-imvision.cn
daiyongya@outlook.com

Abstract

Shaping the behavior of powerful Large Language Models (LLMs) to be both
beneficial and safe is the central challenge of modern AI alignment. We posit
that the post-training alignment process is fundamentally a unified challenge of
Preference Learning, encompassing two distinct modalities: learning from demon-
strated preferences (e.g., Supervised Fine-Tuning, SFT) and from comparative
preferences (e.g., Reinforcement Learning, RL). The current industry-standard
pipeline, which processes these preference types sequentially, is inherently flawed
due to a critical distributional mismatch between the static expert data and the
dynamic policy. This creates two interconnected problems: (1) Offline SFT trains
on a fixed expert distribution, but as the policy's own generation distribution drifts,
the learned knowledge becomes brittle and unreliable. (2) Subsequent online RL
explores to improve generalization, but it operates without direct access to the
rich, ground-truth knowledge within the expert demonstrations, making its explo-
ration inefficient and ungrounded. This fundamental separation prevents the two
data sources from synergistically regularizing each other. To resolve this, we first
reframe alignment as a constrained optimization problem. We then propose Uni-
fied Adversarial Preference Learning (UniAPL), a novel framework that directly
operationalizes this theory by dynamically bridging the gap between the policy's
distribution and the expert's distribution. The ultimate expression of our framework
is a simplified, single-stage unified training objective. This approach cohesively
learns from mixed batches of SFT and preference feedback data, allowing the
dense expert data to directly ground and regularize the online exploration process
in every gradient update. This concurrent optimization inherently mitigates the
distributional mismatch and maximizes data synergy.We empirically validate our
approach on instruction-following tasks using `Qwen3-235B-Instruct-2507` as
the expert teacher. Our model demonstrates comparable or superior general capabil-
ities in English, coding, mathematics, and Chinese, while significantly enhancing
instruction-following ability; it surpasses the strong GRPO baseline by 5.77%
on `Qwen3-0.6B`—matching a 32B model's performance—and exceeds 3.75% on
`Qwen3-4B`, even outperforming the teacher model. Furthermore, analysis of re-
sponse length and log-probability (logp) distributions shows that models trained
with UniAPL not only achieve stronger performance but also generate outputs
closely resembling expert demonstrations.
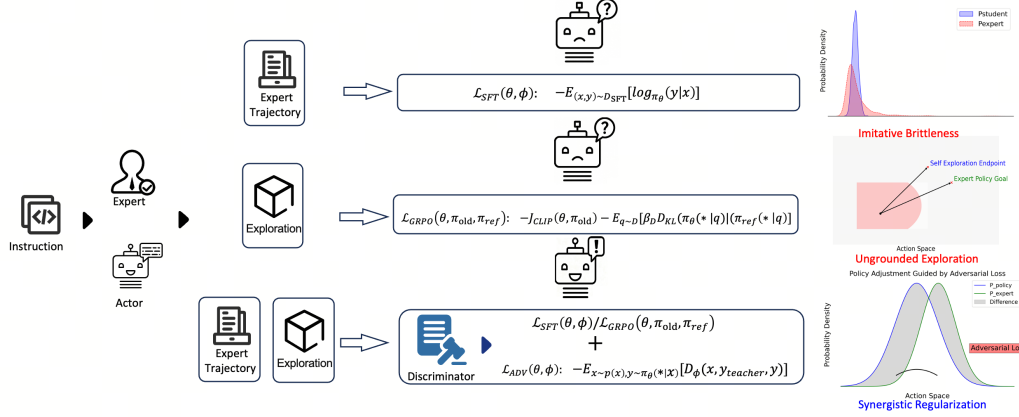
---

*Corresponding author

Figure 1: Overview of the UniAPL training framework. UniAPL unifies offline supervised data and online preference signals into a single, coherent optimization process.

Feature	Non-Adversarial			Adversarial(Ours)		
	SFT	GRPO	SFT→GRPO	A-SFT	A-GRPO	UniAPL
Offline Question	✓	✓	✓	✓	✓	✓
Offline Response	✓	✗	✓	✓	✓	✓
Online Explore	✗	✓	✓	✓	✓	✓
Reward Constraints	✗	✓	✓	✗	✓	✓
Token Level	✓	✓	✓	✓	✓	✓
Sentence Level	✗	✗	✗	✓	✓	✓

Table 1: Illustration of the differences across training paradigms in data usage, exploration strategies, reward formulation, and loss granularity.

1 Introduction

Large Language Models (LLMs) represent a paradigm shift in artificial intelligence, demonstrating remarkable capabilities in complex reasoning, content generation, and human-computer interaction [1, 2]. This transformative potential, however, presents a double-edged sword. As these models grow in power and autonomy, ensuring their behavior remains aligned with human values and intent is not merely a technical refinement but the critical frontier of AI system [3]. The pursuit of Artificial General Intelligence (AGI) hinges on our ability to create systems that are not only intelligent but also helpful, harmless, and honest. Post-training alignment has thus emerged as a foundational field of research, dedicated to shaping these powerful models into reliable partners in human progress [4].

In this work, we introduce a more principled perspective on this challenge: post-training alignment is a unified process of **Preference Learning**. Human intent is conveyed to models through two primary channels of preference data. The first is **Demonstrated Preferences**, where experts provide high-quality examples of desired outputs. Supervised Fine-Tuning (SFT) is the mechanism for learning from this data, instilling the model with foundational knowledge, style, and safety protocols—what we term *Imitative Information* [5]. The second is **Comparative Preferences**, where humans rate or rank different model outputs [6]. Reinforcement Learning (RL) techniques like Direct Preference Optimization (DPO) [7] and Group Relative Policy Optimization (GRPO) [8] are used to learn from this feedback, refining the model's helpfulness and harmlessness.

The fundamental flaw in current methodologies is the **sequential application** of these learning processes, which creates a critical **distributional mismatch** between the static expert data and the dynamic, evolving policy. This mismatch manifests as a severe dilemma with two interconnected problems:

- **The Offline Learning Problem (Imitative Brittleness):** SFT performs offline learning on a fixed, static expert data distribution (P_{expert}). While this provides a crucial knowledge base, it risks overfitting [9]. As the policy begins to generate its own responses, its underlying distribution

2

(P_{policy}) inevitably diverges from the expert's [10]. Due to this gap, the knowledge learned via SFT becomes brittle and fails to generalize reliably to out-of-distribution scenarios.

- **The Online Learning Problem (Ungrounded Exploration):** Subsequent online RL aims to improve generalization by exploring beyond the static expert data. However, this exploration is ungrounded. The policy operates based on its own distribution (P_{policy}) but lacks real-time, direct access to the dense, ground-truth knowledge contained within the expert demonstrations. This makes its exploration inefficient, prone to catastrophic forgetting, and susceptible to unconstrained policy drift towards reward-hacking solutions [11, 12].

This fundamental separation of offline grounding and online exploration prevents the two data sources from synergistically regularizing each other, trapping the alignment process in a state of inefficiency and instability. To resolve this dilemma, we introduce **Unified Adversarial Preference Learning (UniAPL)**, a new paradigm that directly operationalizes a more robust theoretical foundation where alignment is treated as a single, constrained optimization problem. We summarize the key differences in how our method utilizes data, rewards, exploration and loss granularity in Table 1. Our primary contributions are summarized as follows:

- **We re-conceptualize the core challenge of alignment** from simple information decay to a more fundamental problem of distributional mismatch between the brittle, offline-trained policy and the ungrounded, online-exploring policy.
- **We propose a novel framework, UniAPL**, which utilizes a distributional adversarial objective as a dynamic bridge to close this gap, enabling synergistic regularization between demonstrated and comparative preference learning, as illustrated in Figure 1.
- **We introduce a practical, single-stage training paradigm** that cohesively optimizes a unified objective on mixed-data batches. This approach inherently prevents ungrounded drift, maximizes data synergy, and simplifies the entire alignment workflow.
- **We provide extensive empirical validation** for our framework, demonstrating that our unified approach achieves state-of-the-art performance. In particular, we observe absolute improvements of up to 5.77% on `Qwen3-0.6B` and 3.75% on `Qwen3-4B` over strong GRPO baselines.

2 Related Works

Supervised Fine-Tuning has become the cornerstone of post-training for LLMs due to its simplicity and effectiveness. However, by minimizing cross-entropy loss through maximum likelihood estimation, SFT tends to overfit training data, resulting in overconfidence and degraded generalization. To address these limitations, researchers have explored regularization strategies [13, 14, 15] and high-quality data synthesis methods [16, 17, 18, 19, 20] to reduce excessive memorization and knowledge forgetting. More recently, reinforcement learning (RL) has been integrated into the post-training pipeline for its exploratory advantages, and this combination has evolved into the dominant "SFT-then-RL" paradigm, which is now widely adopted in training state-of-the-art LLMs [21, 22, 23, 24, 25] to balance core capability grounding with enhanced adaptability and robustness.

Reinforcement Learning plays a critical role in Post-Training by aligning model outputs with human preferences. Recent work has applied reinforcement learning to domains with verifiable solutions, such as mathematics and programming, leading to reinforcement learning with verifiable rewards (RLVR) and notable performance gains [26, 27, 28]. However, the effectiveness of RL is limited by the model's reliance solely on its own output distribution, without access to external knowledge. On challenging instances, models struggle to produce correct answers, and repeated sampling yields little benefit, leading to suboptimal performance. While approaches such as curriculum learning [29] and dynamic sampling [30] have been proposed to mitigate these issues, their effectiveness remains limited. Recently, incorporating external expert knowledge provides a more straightforward and effective ways to address these difficult cases.

Incorporating external knowledge into reinforcement learning has been shown to improve training efficiency [31, 32, 33]. In the exploration phase of RL, repeated failures on difficult problems indicate that they exceed the model's current ability. In such cases, external guidance can mitigate inefficient exploration on hard problems. Previous studies have leveraged high-quality external responses, either by incorporating them into the rollout process [34] or by providing them as exemplars to guide the policy model [35, 36]. However, the former breaks the on-policy nature of RL, while the latter introduces state inconsistencies between exploration and learning process, leading to suboptimal performance. Some studies [37, 33] introduce cross-entropy loss into RL, but token-level likelihood maximization reduces entropy and ultimately hampers RL effectiveness. Unlike prior approaches,

we employ an adversarial framework where the discriminator enforces indistinguishability between policy outputs and expert demonstrations, and we employ the adversarial framework where the discriminator provides a dynamic, gradient-based regularization signal that guides the student policy towards the semantic manifold of the expert distribution.

3 Methodology: A Unified Framework for Alignment

3.1 Theoretical Foundation: The Alignment Duality and its Dilemma

The objective of LLM alignment is to find an optimal policy π_θ that reflects human intent, conveyed through *Imitative Information* from expert demonstrations D_{SFT} and *Preference Information* from feedback D_{PREF}. An optimally aligned policy must jointly satisfy the constraints from both sources.

Lemma 1 (Optimality Condition for Alignment). *An optimal aligned policy $\pi_{aligned}$ is the solution to the following constrained optimization problem:*

$$\pi_{aligned} = \arg \max_{\pi_\theta} \mathbb{E}_{x \sim p(x), y \sim \pi_\theta(\cdot|x)}[R_\psi(y|x)] \quad subject\ to \quad D(\pi_\theta||\pi^*) \le \epsilon \tag{1}$$

where R_ψ is the reward function from Preference Information, and the constraint, where D is a divergence measure such as KL-divergence, forces the policy π_θ to remain faithful to the expert policy π^*. The sequential "SFT-then-RL" pipeline can be viewed as a specific, often sub-optimal, strategy for solving this problem by decoupling the objectives and optimizing them sequentially.

3.2 The UniAPL Framework: An Adversarial Bridge for the Distributional Gap

UniAPL operationalizes the objective in Lemma 1 by introducing a rich, distributional learning signal. This signal is not merely a component of a loss function, but a distinct gradient vector that guides the optimization process.

The Adversarial Gradient Signal. We employ a discriminator, D_ϕ, to enforce distributional consistency. Given a prompt x, a student response $y_s \sim \pi_\theta(\cdot|x)$, and a teacher response $y_t \sim \pi_{\text{teacher}}(\cdot|x)$, the discriminator outputs a similarity score $D_\phi(x, y_t, y_s) \in [-1, 1]$. The student policy is trained to maximize this score. This process is driven by the adversarial loss $\mathcal{L}_{\text{ADV}}$:

$$\mathcal{L}_{\text{ADV}}(\theta, \phi) = -\mathbb{E}_{x \sim p(x), y \sim \pi_\theta(\cdot|x)}[D_\phi(x, y_{\text{teacher}}, y)]$$

This loss function generates a corresponding **adversarial gradient**, $\mathbf{g}_{\text{ADV}} = \nabla_\theta \mathcal{L}_{\text{ADV}}(\theta, \phi)$, which acts as a vector in the parameter space, pulling the student policy towards the teacher's semantic manifold.

3.3 Gradient Formulations of Alignment Components

Our unified framework views each alignment stage through the lens of the gradient it contributes to the overall policy update.

Adversarial Supervised Fine-Tuning (A-SFT).

Objective. To produce an update direction that simultaneously mimics expert data and matches its underlying distribution.

Methodology. Standard SFT is driven by the negative log-likelihood loss, $\mathcal{L}_{\text{SFT}}(\theta) = -\mathbb{E}_{(x,y) \sim D_{\text{SFT}}}[\log \pi_\theta(y|x)]$, which yields an **imitation gradient**, $\mathbf{g}_{\text{SFT}} = \nabla_\theta \mathcal{L}_{\text{SFT}}(\theta)$. A-SFT enhances this by incorporating the adversarial signal. The total gradient for an A-SFT step is a composite vector:

$$\mathbf{g}_{\text{A-SFT}}(\theta, \phi) = \nabla_\theta \mathcal{L}_{\text{A-SFT}}(\theta, \phi) = \mathbf{g}_{\text{SFT}}(\theta) + \lambda_{\text{adv}} \mathbf{g}_{\text{ADV}}(\theta, \phi)$$

Adversarial Group Relative Policy Optimization (A-GRPO).

Objective. To produce an update direction that seeks higher-reward regions while being grounded by the expert's distribution.

Methodology. This stage utilizes an online RL algorithm, GRPO, which is driven by a reward model RM_ψ trained on D_{PREF}. For a given prompt, G outputs are sampled from a behavior policy π_{old}, and their advantages $\hat{A}_i$ are calculated based on normalized rewards. The GRPO update is derived from its PPO-style loss function, $\mathcal{L}_{\text{GRPO}}$, which combines a clipped surrogate objective $\mathcal{J}_{\text{CLIP}}$ and a KL penalty. This loss is defined as:

$$\mathcal{L}_{\text{GRPO}}(\theta, \pi_{\text{old}}, \pi_{\text{ref}}) = -\left(\mathcal{J}_{\text{CLIP}}(\theta, \pi_{\text{old}}) - \mathbb{E}_{q\sim D}[\beta_D D_{\text{KL}}(\pi_\theta(\cdot|q)||\pi_{\text{ref}}(\cdot|q))]\right)$$

This loss yields a **preference-seeking gradient**, $\mathbf{g}_{\text{GRPO}} = \nabla_\theta \mathcal{L}_{\text{GRPO}}$. A-GRPO augments this with the adversarial signal. The total gradient for an A-GRPO step is therefore:

$$\mathbf{g}_{\text{A-GRPO}}(\theta, \phi) = \nabla_\theta \mathcal{L}_{\text{A-GRPO}} = \mathbf{g}_{\text{GRPO}}(\theta) + \lambda_{\text{adv}}\mathbf{g}_{\text{ADV}}(\theta, \phi)$$

3.4 The Unified Training Objective: A Single-Stage Paradigm

The ultimate expression of UniAPL is a single-stage, unified training objective whose power is best understood by analyzing its gradient.

Methodology. We construct mixed batches from D_{SFT} and D_{PREF} and optimize a single loss function:

$$\mathcal{L}_{\text{Unified}}(\theta, \dots) = \alpha \cdot \mathcal{L}_{\text{A-SFT}}(\theta, \phi) + (1-\alpha) \cdot \mathcal{L}_{\text{A-GRPO}}(\theta, \dots) \tag{2}$$

The Unified Gradient. The essence of our framework is revealed in the gradient of this unified loss, which dictates every parameter update. The total gradient, $\mathbf{g}_{\text{Unified}}$, is a weighted sum of the gradients from the A-SFT and A-GRPO components, computed on their respective data slices within the mixed batch:

$$\mathbf{g}_{\text{Unified}} = \nabla_\theta \mathcal{L}_{\text{Unified}} = \alpha \cdot \mathbf{g}_{\text{A-SFT}}|_{\text{SFT}} + (1-\alpha) \cdot \mathbf{g}_{\text{A-GRPO}}|_{\text{PREF}} \tag{3}$$

By substituting the definitions of the component gradients, we can see that the final update vector is a synergistic combination of four fundamental learning signals:

$$\mathbf{g}_{\text{Unified}} = \underbrace{\alpha\nabla_\theta\mathcal{L}_{\text{SFT}}}_{\substack{\text{Imitation Signal}\\\text{(from SFT data)}}} + \underbrace{(1-\alpha)\nabla_\theta\mathcal{L}_{\text{GRPO}}}_{\substack{\text{Preference Signal}\\\text{(from RL data)}}} + \underbrace{\lambda_{\text{adv}}\left(\alpha\nabla_\theta\mathcal{L}_{\text{ADV}}|_{\text{SFT}} + (1-\alpha)\nabla_\theta\mathcal{L}_{\text{ADV}}|_{\text{PREF}}\right)}_{\substack{\text{Global Distributional}\\\text{Regularization}}}$$

This equation mathematically demonstrates that every single parameter update is simultaneously pushed by gradients for imitation, preference-seeking, and distributional grounding. This is the core of our unified and synergistic optimization process.

Advantages of the Unified Paradigm. By architecting the training process around this single, synergistic gradient, our methodology directly mitigates the core challenges of the sequential paradigm, yielding three distinct advantages:

1. **Inherent Prevention of Ungrounded Policy Drift.** The concurrent training on both SFT and preference data ensures the expert distribution is not a forgotten starting point but a live, dynamic regularizer. The policy is constantly anchored to this ground-truth data manifold, effectively preventing the ungrounded and unchecked policy drift that plagues a separate RL stage.

2. **Synergistic Data Utilization for Enhanced Generalization.** This is more than just data efficiency; it is data synergy. The model learns to balance imitation and preference optimization in every gradient update. The RL signal pushes the model to generalize beyond the potentially overfitted SFT data, while the SFT data provides a rich, grounding signal that makes RL updates more stable and sample-efficient.

3. **Simplified and Efficient Training Workflow.** Beyond the theoretical benefits, our unified approach offers significant operational advantages. The complex, multi-stage training, validation, and model hand-off process is replaced by a single, continuous training run. This dramatically simplifies the overall alignment workflow, reducing engineering overhead and potential sources of error.

These benefits, which we will demonstrate empirically in the following sections, establish our unified approach as a more robust, efficient, and conceptually sound paradigm for LLM alignment.
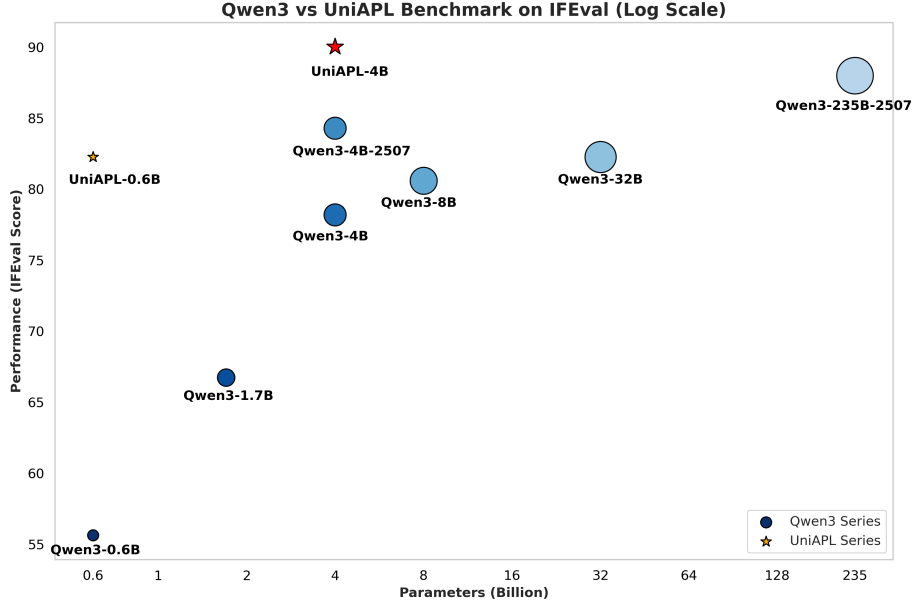
Figure 2: UniAPL performance on the IFEval benchmark. Notably, `UniAPL-0.6B` performs on par with the much larger `Qwen3-32B` model, and `UniAPL-4B` outperforms its teacher model, `Qwen3-235B-A22B-Instruct-2507`, demonstrating the effectiveness and efficiency of the UniAPL approach in enhancing model performance even with significantly smaller model sizes.

4 Experiments

In this section, we evaluate the performance of UniAPL on instruction-following tasks and conduct a comprehensive analysis of the model's output behavior.

4.1 Experimental Setup

Datasets. We use three instruction-following datasets—AutoIF [38], IFevallike [39], and IF-Bench [40]—as our training data, with sizes of 61k, 56k, and 38k examples, respectively. The 61k and 56k examples come from the original AutoIF and IFevallike datasets, while the 38k GT responses in IFBench are obtained by filtering the original 95k examples using a forward pass of a `Qwen2.5-72B` model to retain only those that successfully pass the verification function. Each example consists of a prompt, a verification function, and a corresponding ground-truth response (GT) that passes the verification function.Examples of the training data, validation functions, and response data can be found in the appendix B.4.

Teacher Policy Responses. The teacher policy responses are generated by performing forward inference with the `Qwen3-235B-Instruct-2507` model on the prompts from the datasets. Note that these responses are not guaranteed to pass the verification function.

Benchmarks. To verify the effectiveness of our method,we evaluate it on two common instruction-following benchmarks and ten general-purpose tasks,including (IFEval [41], MultiIF [42]), general English (MMLU [43], MMLU-Pro [44], GPQA-Diamond [45]), coding(Humaneval [46], Mbpp [47]) and mathematics (GSM8K [48], Math-500 [49], TheoremQA [50]), and Chinese (CMMLU [51], CEval [52]). We report the average performance per domain in the main tables, with detailed results for each dataset provided in the Appendix B.3.

Setting. According to our theoretical insight, a greater discrepancy between the teacher and student policies is beneficial for performance;To validate the generality and effectiveness of UniAPL, we incorporate it into both supervised fine-tuning (SFT) and reinforcement learning (RL) settings, training Qwen3 models of two sizes: 0.6B and 4B. In the SFT setting, we directly compute the loss using standard cross-entropy between predictions and ground-truth responses. For RL, we directly adopt GRPO.For the verifiable reward, the model receives a reward of 1 if its response passes the verification function,and 0 otherwise. We evaluate the impact of integrating UniAPL at

Method	Benchmarks					
	Avg	I-Following	English	Coding	Mathematics	Chinese
Qwen3-0.6B	40.81	56.62	32.61	32.63	42.29	43.25
SFT	38.48	59.97	25.20	29.43	39.56	44.33
GRPO	40.93	74.09	32.39	17.58	41.49	43.11
SFT→GRPO	42.56	77.40	35.09	21.60	39.57	44.38
SFT→A-GRPO	43.61	79.57	32.71	29.00	39.95	44.10
A-SFT	39.83	60.72	26.67	**33.93**	40.14	44.10
A-GRPO	41.99	79.58	31.36	19.79	41.04	43.96
A-SFT→A-GRPO	**43.31**	**79.86**	**36.48**	33.72	**43.72**	**44.47**
Qwen3-4B	62.52	80.10	45.42	70.31	59.47	67.36
SFT	60.80	80.60	41.13	68.62	56.92	68.51
GRPO	68.11	86.90	53.42	71.43	**66.52**	70.42
SFT→GRPO	65.75	89.75	49.04	71.54	59.40	70.53
SFT→A-GRPO	68.11	90.45	48.25	72.06	60.53	70.93
A-SFT	60.67	82.47	38.78	70.34	56.65	68.07
A-GRPO	**68.39**	88.17	**54.29**	73.15	65.12	69.89
A-SFT→A-GRPO	67.93	**90.65**	51.16	**73.26**	63.65	**71.45**

Table 2: Comparative results of different training methods across multiple benchmarks (IFEval, MulitIF, MMLU, MMLU-Pro, GPQA-Diamond, HumanEval, MBPP, GSM8K, MATH-500, TheoremQA, CMMLU, CEval) for models with `Qwen3-0.6B` and `Qwen3-4B`.

Method	Benchmarks					
	Avg	I-Following	English	Coding	Mathematics	Chinese
Qwen3-0.6B	40.81	56.62	32.61	32.63	**42.29**	43.25
CHORD	41.19	73.52	29.92	26.09	38.52	44.90
A-CHORD	**42.09**	**75.25**	**32.74**	23.89	39.45	**45.14**
Qwen3-4B	62.52	80.10	45.42	70.31	59.47	67.36
CHORD	**69.30**	86.97	54.86	**73.76**	67.67	71.30
A-CHORD	69.27	**88.27**	**55.46**	70.32	**68.14**	**71.66**

Table 3: Comparative results of unified training methods across multiple benchmarks (IFEval, MulitIF, MMLU, MMLU-Pro, GPQA-Diamond, HumanEval, MBPP, GSM8K, MATH-500, TheoremQA, CMMLU, CEval) for models with `Qwen3-0.6B` and `Qwen3-4B`.

different stages: A-SFT denotes the addition of the UniAPL adversaril loss to the cross-entropy loss, while A-GRPO represents the incorporation of the adversaril loss into the GRPO objective. When computing distances using UniAPL, to avoid introducing a discriminator that would be progressively updated during training, we adopt and modify POLAR [53] as a universal discriminator. POLAR is a reference-based reward model that evaluates the likelihood that a response and its reference originate from the same policy model and outputs a corresponding score. The details of our modifications to adapt POLAR into a universal discriminator are provided in the Appendix A.3. Finally, we use CHORD [33], a method that jointly performs supervised fine-tuning (SFT) and reinforcement learning (RL), as a hybrid training baseline to validate the generality of our adversarial loss. All experiments were based on MS-Swift [54],VeRL [55] and Trinity-Rft [56].Detailed experimental settings can be found in Appendix B.1.

4.2 Main Result

Objective Metrics and Experimental Analysis. Our primary experimental results focus on two training paradigms: staged training and unified training with the UniAPL adversarial loss. As shown in Table 2 and Table 3, several key findings emerge from the staged training setup. First,
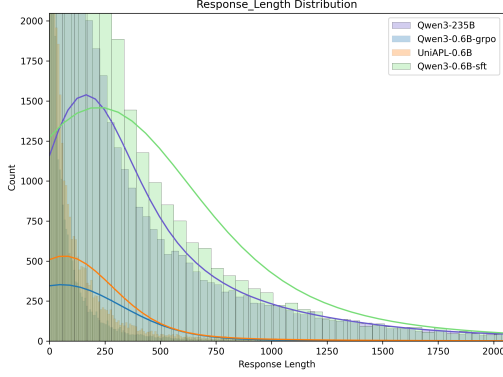
Figure 3: Response length distributions on the IFBench-38K dataset under teacher models and different training paradigms. Among them, SFT is most similar to the teacher, while UniAPL achieves a response length distribution more consistent with the teacher than GRPO.
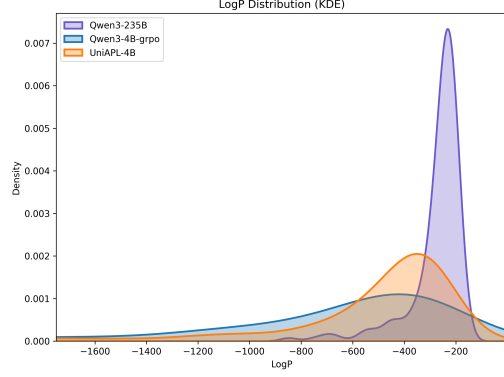
Figure 4: Kernel Density Estimation (KDE) of log-probability differences between student and teacher models trained with GRPO and UniAPL. UniAPL exhibits a narrower distribution and reduced divergence, indicating closer alignment with the teacher.

in the offline supervised fine-tuning (SFT) phase, A-SFT slightly mitigates the decline in general capabilities, achieving a better performance gain on instruction-following tasks compared to standard SFT. Second, during the online training phase, we observe that A-GRPO surprisingly achieves performance comparable to the two-stage SFT→GRPO approach on task-specific metrics. This is likely because A-GRPO continuously guides the student policy to align with the teacher policy during training, effectively serving as an implicit SFT warm-up. This effect is particularly evident in the 4B model, where A-GRPO directly prevents the degradation of general capabilities typically introduced by SFT. Furthermore, in both our two-stage and unified training experiments, the UniAPL adversarial loss still leads to improved performance.

Moreover, by enhancing the model's instruction-following capability, the adversarial loss enables deeper comprehension of instructions, leading to superior performance across various capabilities compared to the base model.

Theoretical Validation. To verify that UniAPL effectively learns from the teacher model while mitigating policy drift, we re-evaluate models trained with SFT, GRPO, and UniAPL on the 38K instances from IFBench. We report the distribution of response lengths in Figure 3. Our results show that SFT captures only the surface-level length pattern, while GRPO relies on self-exploration. UniAPL, by emulating the teacher's reasoning, produces responses whose length distribution aligns more closely with the teacher than GRPO does. An example can be found in Appendix B.5. In addition, to verify that UniAPL continuously constrains the student model to learn from the teacher, we sample 500 responses from the dataset and compute the total log-likelihood (logp) under three models: the teacher, the student trained with GRPO, and the student trained with UniAPL. The results are summarized in Figure 4. The narrower distribution and reduced divergence under UniAPL indicate closer alignment between student and teacher outputs, demonstrating its superior policy regularization and more effective knowledge transfer compared to GRPO.

4.3 Ablation Study

Ablation on UniAPL Adversaril Loss Method. We validate the effectiveness of the base model and ground-truth (GT) Dataset in the appendix B.1. Furthermore, we analyze the impact of the UniAPL discriminator coefficients in conjunction with GT references and the KL divergence loss in Appendix B.2. To investigate the contribution of the UniAPL adversaril loss,we conducted a series of ablation experiments under the SFT→AGRPO training paradigm using the base model and GT Dataset. Specifically, we compared three settings in Table 4, where the base method employs RLVR (Reinforcement Learning with Verifiable Rewards)'s GRPO. (i) directly adding the raw output of POLAR to the verifiable reward, (ii) adding the transformed discriminator output to the verifiable reward, and (iii) using the transformed discriminator output as a separate adversaril loss term. Among these settings, using it as an independent adversaril loss achieved the best performance,which aligns

Method	Benchmarks					
	Avg	I-Following	English	Coding	Mathematics	Chinese
RLVR	47.53	76.25	**36.43**	36.29	43.83	52.22
RLVR+POLAR	46.70	68.47	36.01	36.80	**46.06**	51.82
RLVR+A_{coef}	47.68	73.09	35.96	**38.94**	45.09	**52.51**
w A_{coef}	**47.83**	**78.94**	36.23	35.57	43.91	52.26

Table 4: RLVR denotes using only verifiable rewards, corresponding to the SFT+GRPO setup. RLVR+POLAR represents setting (i), while RLVR+ADV$_{\text{coef}}$ indicates directly incorporating the reward model's output into the original GRPO reward, corresponding to setting (ii), where ADV$_{\text{coef}}$ refers to the discriminator's output. Finally, w ADV$_{\text{coef}}$ denotes the proposed method.
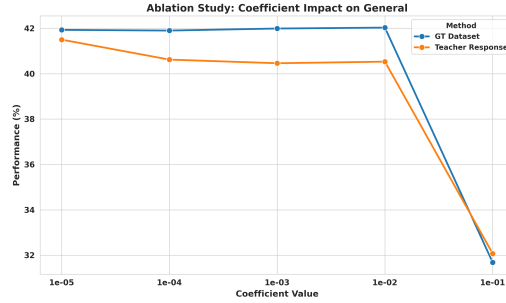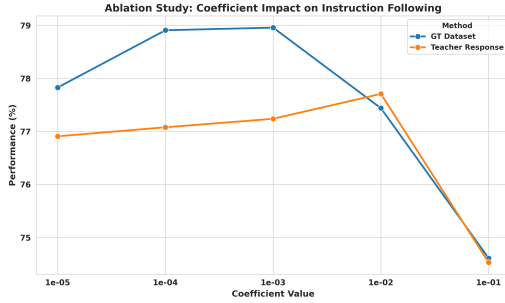

Figure 5: Ablation study of discriminator coefficients in UniAPL on instruction-following tasks.

Figure 6: Ablation study of discriminator coefficients in UniAPL on general performance.

with our theoretical expectation. In contrast, directly adding it to the verifiable reward tended to overwhelm the original reward signal.

Ablation on the UniAPL Adversaril Loss Coefficient λ_{adv}. We conducted extensive ablation studies on the UniAPL adversarial loss coefficient λ_{adv}, varying it from 0.1 to 1×10^{-5}, which spans the typical magnitude difference relative to the policy gradient loss (PGloss). In all experiments, the coefficient of the KL loss was fixed at 0.001. Figure 5 and 6 presents all coefficient settings evaluated on both GT-based data and teacher-generated responses. We suggest setting the coefficient within the range of 0.01 to 0.0001, which is approximately one order of magnitude lower than the policy gradient loss, serving as an auxiliary signal. A value that is too high may dominate the optimization with the adversarial term and suppress the model's exploration, while a value that is too low may lead to insufficient performance gains.

5 Conclusion

This work introduces Unified Adversarial Preference Learning (UniAPL), a principled framework designed to mitigate the distributional mismatch in sequential SFT-then-RL alignment. Our central conceptual contribution is the reframing of alignment as a single-stage, constrained optimization problem, rather than a sequence of disparate steps. The framework realizes this through an elegant yet effective adversarial objective that enforces distributional consistency, allowing dense SFT signals to dynamically regularize the online exploration of RL. This concurrent process fosters a crucial synergy: expert data grounds exploration to prevent policy drift, while preference learning pushes the model to generalize beyond the static expert distribution. While our approach is primarily methodological, its conceptual soundness is validated by significant empirical gains and a streamlined training workflow. These results substantiate our core hypothesis that unifying imitative and preference-based learning signals is a more effective and efficient path to alignment. Looking forward, UniAPL serves as a practical and extensible paradigm, offering a foundation for integrating more diverse supervisory signals and advancing toward more robust and scalable methods for steering powerful AI systems.

References

[1] Jingfeng Yang, Hongye Jin, Ruixiang Tang, Xiaotian Han, Qizhang Feng, Haoming Jiang, Shaochen Zhong, Bing Yin, and Xia Hu. Harnessing the power of llms in practice: A survey on chatgpt and beyond. ACM Transactions on Knowledge Discovery from Data, 18(6):1–32, 2024.

[2] Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, et al. Sparks of artificial general intelligence: early experiments with gpt-4 (2023). arXiv preprint arXiv:2303.12712, 1, 2023.

[3] Youbang Sun, Xiang Wang, Jie Fu, Chaochao Lu, and Bowen Zhou. R^2AI: Towards resistant and resilient ai in an evolving world. arXiv preprint arXiv:2509.06786, 2025.

[4] Guiyao Tie, Zeli Zhao, Dingjie Song, Fuyang Wei, Rong Zhou, Yurou Dai, Wen Yin, Zhejian Yang, Jiangyue Yan, Yao Su, et al. A survey on post-training of large language models. arXiv e-prints, pages arXiv–2503, 2025.

[5] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. Advances in neural information processing systems, 35:27730–27744, 2022.

[6] Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. Advances in neural information processing systems, 30, 2017.

[7] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. Advances in neural information processing systems, 36:53728–53741, 2023.

[8] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Yang Wu, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. arXiv preprint arXiv:2402.03300, 2024.

[9] Minghuan Liu, Hanye Zhao, Zhengyu Yang, Jian Shen, Weinan Zhang, Li Zhao, and Tie-Yan Liu. Curriculum offline imitating learning. Advances in Neural Information Processing Systems, 34:6266–6277, 2021.

[10] Lantao Yu, Tianhe Yu, Jiaming Song, Willie Neiswanger, and Stefano Ermon. Offline imitation learning with suboptimal demonstrations via relaxed distribution matching. In Proceedings of the AAAI conference on artificial intelligence, volume 37, pages 11016–11024, 2023.

[11] Huaijie Wang, Shibo Hao, Hanze Dong, Shenao Zhang, Yilin Bao, Ziran Yang, and Yi Wu. Offline reinforcement learning for llm multi-step reasoning. arXiv preprint arXiv:2412.16145, 2024.

[12] Thomas Carta, Clément Romac, Thomas Wolf, Sylvain Lamprier, Olivier Sigaud, and Pierre-Yves Oudeyer. Grounding large language models in interactive environments with online reinforcement learning. In International Conference on Machine Learning, pages 3676–3713. PMLR, 2023.

[13] Zhihan Liu, Miao Lu, Shenao Zhang, Boyi Liu, Hongyi Guo, Yingxiang Yang, Jose Blanchet, and Zhaoran Wang. Provably mitigating overoptimization in rlhf: Your sft loss is implicitly an adversarial regularizer. Advances in Neural Information Processing Systems, 37:138663–138697, 2024.

[14] Ziniu Li, Congliang Chen, Tian Xu, Zeyu Qin, Jiancong Xiao, Ruoyu Sun, and Zhi-Quan Luo. Entropic distribution matching in supervised fine-tuning of llms: Less overfitting and better diversity.(2024). URL https://arxiv. org/abs/2408.16673.

[15] Gabriel Pereyra, George Tucker, Jan Chorowski, Łukasz Kaiser, and Geoffrey Hinton. Regularizing neural networks by penalizing confident output distributions. arXiv preprint arXiv:1701.06548, 2017.

[16] Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B Hashimoto. Stanford alpaca: An instruction-following llama model, 2023.

[17] Ganqu Cui, Lifan Yuan, Ning Ding, Guanming Yao, Bingxiang He, Wei Zhu, Yuan Ni, Guotong Xie, Ruobing Xie, Yankai Lin, et al. Ultrafeedback: Boosting language models with scaled ai feedback. arXiv preprint arXiv:2310.01377, 2023.

[18] Hao Chen, Abdul Waheed, Xiang Li, Yidong Wang, Jindong Wang, Bhiksha Raj, and Marah I Abdin. On the diversity of synthetic data and its impact on training large language models, 2024. URL https://arxiv. org/abs/2410.15226, 2016.

[19] Maosongcao Maosongcao, Taolin Zhang, Mo Li, Chuyu Zhang, Yunxin Liu, Conghui He, Haodong Duan, Songyang Zhang, and Kai Chen. Condor: Enhance llm alignment with knowledge-driven data synthesis and refinement. In Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pages 22392–22412, 2025.

[20] Fan Lin, Shuyi Xie, Yong Dai, Wenlin Yao, Tianjiao Lang, and Yu Zhang. Idgen: Item discrimination induced prompt generation for llm evaluation. Advances in Neural Information Processing Systems, 37:88557–88580, 2024.

[21] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. arXiv preprint arXiv:2505.09388, 2025.

[22] Aixin Liu, Bei Feng, Bing Xue, Bingxuan Wang, Bochao Wu, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, et al. Deepseek-v3 technical report. arXiv preprint arXiv:2412.19437, 2024.

[23] Gemma Team, Thomas Mesnard, Cassidy Hardin, Robert Dadashi, Surya Bhupatiraju, Shreya Pathak, Laurent Sifre, Morgane Rivière, Mihir Sanjay Kale, Juliette Love, et al. Gemma: Open models based on gemini research and technology. arXiv preprint arXiv:2403.08295, 2024.

[24] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. arXiv preprint arXiv:2303.08774, 2023.

[25] Pengyu Cheng, Yifan Yang, Jian Li, Yong Dai, Tianhao Hu, Peixin Cao, Nan Du, and Xiaolong Li. Adversarial preference optimization: Enhancing your alignment via rm-llm game. arXiv preprint arXiv:2311.08045, 2023.

[26] Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, et al. Tulu 3: Pushing frontiers in open language model post-training. arXiv preprint arXiv:2411.15124, 2024.

[27] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. arXiv preprint arXiv:2501.12948, 2025.

[28] Youssef Mroueh. Reinforcement learning with verifiable rewards: Grpo's effective loss, dynamics, and success amplification. arXiv preprint arXiv:2503.06639, 2025.

[29] Xiao Liu, Bo Qin, Dongzhu Liang, Guang Dong, Hanyu Lai, Hanchen Zhang, Hanlin Zhao, Iat Long Iong, Jiadai Sun, Jiaqi Wang, et al. Autoglm: Autonomous foundation agents for guis. arXiv preprint arXiv:2411.00820, 2024.

[30] Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Tiantian Fan, Gaohong Liu, Lingjun Liu, Xin Liu, et al. Dapo: An open-source llm reinforcement learning system at scale, 2025. URL https://arxiv. org/abs/2503.14476, 2025.

[31] Tianzhe Chu, Yuexiang Zhai, Jihan Yang, Shengbang Tong, Saining Xie, Dale Schuurmans, Quoc V Le, Sergey Levine, and Yi Ma. Sft memorizes, rl generalizes: A comparative study of foundation model post-training. arXiv preprint arXiv:2501.17161, 2025.

[32] Joey Hong, Anca Dragan, and Sergey Levine. Q-sft: Q-learning for language models via supervised fine-tuning. arXiv preprint arXiv:2411.05193, 2024.

[33] Wenhao Zhang, Yuexiang Xie, Yuchang Sun, Yanxi Chen, Guoyin Wang, Yaliang Li, Bolin Ding, and Jingren Zhou. On-policy rl meets off-policy experts: Harmonizing supervised fine-tuning and reinforcement learning via dynamic weighting. arXiv preprint arXiv:2508.11408, 2025.

[34] Jianhao Yan, Yafu Li, Zican Hu, Zhi Wang, Ganqu Cui, Xiaoye Qu, Yu Cheng, and Yue Zhang. Learning to reason under off-policy guidance, 2025. URL https://arxiv. org/abs/2504.14945.

11

[35] Jianhao Yan, Yafu Li, Zican Hu, Zhi Wang, Ganqu Cui, Xiaoye Qu, Yu Cheng, and Yue Zhang. Learning to reason under off-policy guidance. arXiv preprint arXiv:2504.14945, 2025.

[36] Xuechen Zhang, Zijian Huang, Yingcong Li, Chenshun Ni, Jiasi Chen, and Samet Oymak. Bread: Branched rollouts from expert anchors bridge sft & rl for reasoning. arXiv preprint arXiv:2506.17211, 2025.

[37] Yuqian Fu, Tinghong Chen, Jiajun Chai, Xihuai Wang, Songjun Tu, Guojun Yin, Wei Lin, Qichao Zhang, Yuanheng Zhu, and Dongbin Zhao. Srft: A single-stage method with supervised and reinforcement fine-tuning for reasoning. arXiv preprint arXiv:2506.19767, 2025.

[38] Guanting Dong, Keming Lu, Chengpeng Li, Tingyu Xia, Bowen Yu, Chang Zhou, and Jingren Zhou. Self-play with execution feedback: Improving instruction-following capabilities of large language models. arXiv preprint arXiv:2406.13542, 2024.

[39] Zhangchen Xu, Fengqing Jiang, Luyao Niu, Yuntian Deng, Radha Poovendran, Yejin Choi, and Bill Yuchen Lin. Magpie: Alignment data synthesis from scratch by prompting aligned llms with nothing, 2024.

[40] Valentina Pyatkin, Saumya Malik, Victoria Graf, Hamish Ivison, Shengyi Huang, Pradeep Dasigi, Nathan Lambert, and Hannaneh Hajishirzi. Generalizing verifiable instruction following, 2025.

[41] Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou. Instruction-following evaluation for large language models. arXiv preprint arXiv:2311.07911, 2023.

[42] Yun He, Di Jin, Chaoqi Wang, Chloe Bi, Karishma Mandyam, Hejia Zhang, Chen Zhu, Ning Li, Tengyu Xu, Hongjiang Lv, et al. Multi-if: Benchmarking llms on multi-turn and multilingual instructions following. arXiv preprint arXiv:2410.15553, 2024.

[43] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. arXiv preprint arXiv:2009.03300, 2020.

[44] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyan Jiang, et al. Mmlu-pro: A more robust and challenging multi-task language understanding benchmark. Advances in Neural Information Processing Systems, 37:95266–95290, 2024.

[45] David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R Bowman. Gpqa: A graduate-level google-proof q&a benchmark. In First Conference on Language Modeling, 2024.

[46] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde De Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. arXiv preprint arXiv:2107.03374, 2021.

[47] Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. arXiv preprint arXiv:2108.07732, 2021.

[48] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. arXiv preprint arXiv:2110.14168, 2021.

[49] Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let's verify step by step. In The Twelfth International Conference on Learning Representations, 2023.

[50] Wenhu Chen, Ming Yin, Max Ku, Pan Lu, Yixin Wan, Xueguang Ma, Jianyu Xu, Xinyi Wang, and Tony Xia. Theoremqa: A theorem-driven question answering dataset. In The 2023 Conference on Empirical Methods in Natural Language Processing, 2023.

[51] Haonan Li, Yixuan Zhang, Fajri Koto, Yifei Yang, Hai Zhao, Yeyun Gong, Nan Duan, and Timothy Baldwin. Cmmlu: Measuring massive multitask language understanding in chinese, 2023.

[52] Yuzhen Huang, Yuzhuo Bai, Zhihao Zhu, Junlei Zhang, Jinghan Zhang, Tangjun Su, Junteng Liu, Chuancheng Lv, Yikai Zhang, Jiayi Lei, Yao Fu, Maosong Sun, and Junxian He. C-eval:

A multi-level multi-discipline chinese evaluation suite for foundation models. In <u>Advances in Neural Information Processing Systems</u>, 2023.

[53] Shihan Dou, Shichun Liu, Yuming Yang, Yicheng Zou, Yunhua Zhou, Shuhao Xing, Chenhao Huang, Qiming Ge, Demin Song, Haijun Lv, et al. Pre-trained policy discriminators are general reward models. <u>arXiv preprint arXiv:2507.05197</u>, 2025.

[54] Yuze Zhao, Jintao Huang, Jinghan Hu, Xingjun Wang, Yunlin Mao, Daoze Zhang, Hong Zhang, Zeyinzi Jiang, Zhikai Wu, Baole Ai, Ang Wang, Wenmeng Zhou, and Yingda Chen. Swift:a scalable lightweight infrastructure for fine-tuning, 2025.

[55] Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. Hybridflow: A flexible and efficient rlhf framework. In <u>Proceedings of the Twentieth European Conference on Computer Systems</u>, EuroSys '25, page 1279–1297. ACM, March 2025.

[56] Xuchen Pan, Yanxi Chen, Yushuo Chen, Yuchang Sun, Daoyuan Chen, Wenhao Zhang, Yuexiang Xie, Yilun Huang, Yilei Zhang, Dawei Gao, Yaliang Li, Bolin Ding, and Jingren Zhou. Trinity-rft: A general-purpose and unified framework for reinforcement fine-tuning of large language models, 2025.

A Theoretical and Implementation Details of the UniAPL Framework

This appendix provides a deeper exploration of the theoretical underpinnings and implementation details of our proposed framework.

A.1 The Unified View: SFT as a Special Case of Preference Learning

Our framework is built on the premise that SFT and RL are not disparate stages but points on a continuum of preference learning. To formalize this, we re-interpret SFT as an extreme and limiting case of preference learning.

The Conventional View: SFT as Imitation Learning. Traditionally, SFT is understood as imitation learning. The objective is to train a policy π_θ to mimic an expert by minimizing the Negative Log-Likelihood (NLL) of a single ground-truth response y^*:

$$\mathcal{L}_{\text{SFT}}(\theta) = -\log \pi_\theta(y^*|x)$$

Minimizing this loss is equivalent to maximizing the probability $\pi_\theta(y^*|x)$. In this view, SFT is about replicating a single correct answer, distinct from preference learning, which involves choosing between multiple answers.

The Unified View: SFT as Preference Learning with a Dirac Delta Reward. Our framework unifies these concepts by redefining "preference" for the SFT case. While standard preference learning (e.g., DPO) is relative (a "winner" y_w is preferred over a "loser" y_l), SFT can be seen as expressing an absolute and infinite preference.

Specifically, the ground-truth response y^* is **infinitely preferred** over any other possible response $y \neq y^*$. There is no concept of a response being "close" or "better" than another incorrect response. To model this, we define a reward function using the Dirac delta function, $\delta(\cdot)$:

$$R(y|x) = \delta(y - y^*)$$

This function grants an infinite reward if the policy generates the exact ground-truth response and zero reward for any deviation.

Mathematical Equivalence of Reward Maximization and NLL Minimization. We can formally show that the standard SFT objective is a special case of the general reward-maximization objective.

1. The standard reinforcement learning objective is to maximize the expected reward for a response y sampled from the policy $\pi_\theta(\cdot|x)$:

$$\max_\theta \mathbb{E}_{y \sim \pi_\theta(\cdot|x)}[R(y|x)]$$

2. Substituting our Dirac delta reward function:

$$\max_\theta \mathbb{E}_{y \sim \pi_\theta(\cdot|x)}[\delta(y - y^*)]$$

3. The expectation is an integral (or sum) of the reward of each possible output y weighted by its probability $\pi_\theta(y|x)$. Due to the properties of the Dirac delta function, the integral is non-zero only at the single point $y = y^*$. Therefore, the expectation collapses to be proportional to the probability of generating y^*:

$$\mathbb{E}_{y \sim \pi_\theta(\cdot|x)}[\delta(y - y^*)] \propto \pi_\theta(y^*|x)$$

4. To maximize this expectation, one must maximize the probability term:

$$\max_\theta \pi_\theta(y^*|x)$$

5. Since the logarithm is a monotonic function, this is equivalent to maximizing its logarithm:

$$\max_\theta \log \pi_\theta(y^*|x)$$

6. Finally, maximizing a function is equivalent to minimizing its negation:

$$\min_\theta(-\log \pi_\theta(y^*|x))$$

This final expression is precisely the NLL loss for SFT. We have formally shown that the objective of SFT is a special case of the general reward-maximization objective, where the reward function is a Dirac delta centered on the expert demonstration.

A.2 Theoretical Significance of the Unified View

This unification is critically important for several reasons:

1. **Establishes a Coherent Theoretical Framework.** It places SFT, DPO, and online RL under a single, elegant objective function, $J(\theta)$. This provides a cohesive narrative for the entire post-training pipeline, viewing it not as a sequence of disparate steps but as a unified process.

2. **Reveals a Curriculum of Reward Signals.** It frames post-training as a principled progression where the nature of the reward signal evolves. Stage 1 (SFT) uses a sharp, sparse Dirac delta signal to instill absolute knowledge. Subsequent stages (DPO/GRPO) replace this sharp signal with a smooth preference landscape, teaching the model to navigate nuanced trade-offs and generalize in a continuous space of response quality.

A.3 Implementation of the Adversarial Discriminator

As detailed in the main text, our adversarial objective $\mathcal{L}_{\text{ADV}}$ relies on a discriminator D_ϕ. To implement this, we adapt the methodology from POLAR, a model designed to measure the distance between generation strategies. We use a pre-trained POLAR model as the backbone for our discriminator. The raw distance score from POLAR is then transformed into a well-behaved adversarial loss coefficient for our objective. The full procedure is detailed in Algorithm 1.

Algorithm 1 Adversarial Loss Coefficient Computation Based on POLAR Distance

---

Require: Prompt x, Teacher policy π_t, Student policy π_s
Ensure: Adversarial loss coefficient $coef \in [-1, 1]$
 1: Obtain teacher response $y_t = \pi_t(x)$ and student response $y_s = \pi_s(x)$
 2: Feed (x, y_t, y_s) into POLAR model to get BT distance: $r = \text{POLAR}(x, y_t, y_s)$
 3: Normalize distance r into similarity score p using sigmoid:

$$p = \sigma(r) = \frac{1}{1 + e^{-r}} \quad (p \in (0, 1))$$

 4: Convert similarity score p into an adversarial loss coefficient $coef$:

$$coef = 1 - 8(p - 0.5)^2 \quad (coef \in [-1, 1])$$

 5: **return** $coef$

The final $coef$ serves as the output of our discriminator, $D_\phi(x, y_t, y_s)$, which is then used to compute $\mathcal{L}_{\text{ADV}}$.

A.4 Alternative Preference Optimization Formulations

Our UniAPL framework is compatible with a wide range of preference optimization algorithms. While the main text focuses on a policy-gradient approach (GRPO), here we detail prominent reward-model-free alternatives.

A.4.1 Direct Preference Optimization (DPO)

DPO is a widely-used algorithm that formulates preference learning as a direct classification problem on human preferences.

Data Format. The preference dataset D_{PREF} for DPO consists of tuples (x, y_w, y_l), where for a given prompt x, y_w is the "winner" (preferred) response and y_l is the "loser" (dispreferred) response.

Loss Function. The DPO loss is formulated to directly increase the likelihood of the winner response while decreasing the likelihood of the loser response, relative to a fixed reference policy π_{ref}. It is defined as:

$$\mathcal{L}_{\text{DPO}}(\theta, \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim D_{\text{PREF}}} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi_\theta(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right) \right]$$

Here, β is a temperature parameter that controls the strength of the preference modeling, and σ is the sigmoid function.

A.4.2 Kahneman-Tversky Optimization (KTO)

KTO offers a different perspective by removing the need for pairwise comparisons altogether. It instead relies on binary labels for individual examples.

Data Format. The preference dataset D_{PREF} for KTO consists of tuples (x, y, l), where $l \in \{\text{desirable}, \text{undesirable}\}$ is a binary label indicating whether the response y to prompt x is good or bad.

Objective. The KTO loss is designed based on principles from human prospect theory. It has two main components: one term encourages the policy to increase the likelihood of "desirable" examples, and another, more heavily weighted term, strongly discourages the policy from generating "undesirable" examples. This asymmetry reflects the human tendency to be more sensitive to losses than to equivalent gains.

B Detailed benchmark results And Example

B.1 Detailed experiment training setting

General setting All experiments were conducted on eight H100 GPUs.The learning rate was set to 1×10^{-5} for the SFT stage, 1×10^{-7} for the GRPO stage, and 1×10^{-6} for the mixed training following the default configuration of CHORD.

Chat model setting Since the Qwen3 models require manual control of the *think* process, all GRPO experiments based on the Chat models were performed with `enable_thinking=False`. Each model response begins with the prefix `<think>\n\n</think>\n\n`. Similarly, for SFT training on Chat models, this prefix was prepended to the response part of the data,and the loss on empty think segments was excluded from the computation.

Base model Setting To validate the effectiveness of our approach on the base model and ground-truth references, we conducted several ablation studies on the Qwen3-Base model using the reference answers. To eliminate the potential interference of reasoning-related labels, we applied supervised fine-tuning (SFT) with the `Qwen2.5` chat template, and removed the `<think>\n\n</think>\n\n` prefix from the SFT response data. Since A-SFT requires an additional rollout step and the base model lacks inherent conversational ability, experiments based on the base model were carried out under the SFT→GRPO training paradigm. As shown in Table 5, when using the ground-truth (GT) reference responses from the dataset as evaluation targets, initializing from the base model achieves better generalization performance than initializing from a chat model, while reaching comparable instruction-following capability.

Method	Benchmarks					
	Avg	I-Following	English	Coding	Mathematics	Chinese
0.6B-Basenothink/GT	-	-	-	-	-	-
SFT	43.97	60.13	33.32	35.87	43.60	**52.43**
SFT→ GRPO	47.53	76.25	**36.43**	**36.29**	43.83	52.22
SFT→ A-GRPO	**47.83**	**78.96**	36.23	35.57	**43.91**	52.26

Table 5: Experimental results of base model on the ground-truth (GT) dataset.SFT→GRPO corresponds to RLVR in Table 12, while SFT and SFT→A-GRPO corresponds to the entry with rate 0.001 in Table 10.

B.2 Supplementary ablation analysis

Ablation on UniAPL discriminator coefficients with GT References and KL Loss We further conducted an ablation study on the UniAPL discriminator coefficient using ground-truth (GT) reference answers.Specifically, the GT answers were treated as the target, and both the teacher and student outputs were passed through POLAR to obtain their respective scores, r_t and r_s. In the policy space, this corresponds to the teacher and student policies lying within a circle centered at the reference policy with radius $\max(r_t, r_s)$. The difference between these scores, Δr, approximates the distance R between the two policies. Subsequently, the distances between these scores were calculated according to the algorithm described in the Appendix A.3 and used to derive the discriminator coefficient. Table 6 presents the results of this ablation study. Finally, we conducted ablation experiments on the KL loss to examine its effect on training stability and performance. When the KL loss was disabled, the model's policy optimization was guided solely by the PG loss and adversaril loss. Through this experiment,the presence or absence of reference answers in the calculation of the discriminator coefficient has little impact on the final performance.Disabling the KL loss allows the model to achieve better performance on this specific task,but it may leads to a collapse of its general capabilities on other tasks.

Method	Benchmarks					
	Avg	I-Following	English	Coding	Mathematics	Chinese
Directly Distinguish	-	-	-	-	-	-
Use KL	45.86	77.24	33.29	**37.41**	38.84	**52.31**
NO KL	**47.52**	**80.65**	**36.03**	35.39	**41.90**	52.21
With GT Response	-	-	-	-	-	-
Use KL	**46.25**	77.43	32.48	36.40	41.98	**51.95**
NO KL	40.65	**80.38**	**36.62**	**38.01**	12.64	51.62

Table 6: Ablation study of UniAPL discriminator coefficients with ground-truth references and KL loss.

B.3 Detailed benchmark results

Benchmark	0.6B	SFT	GRPO	S→G	S→A-G	A-SFT	A-GRPO	A-S→A-G
Avg	40.81	38.48	40.93	42.56	43.61	39.83	41.99	43.31
IFEVAL	55.64	60.26	74.68	80.78	82.07	61.74	82.26	82.99
MultiIF	57.60	59.68	73.50	74.01	77.06	59.70	76.89	76.73
Mmlu	43.91	46.16	42.01	43.79	43.35	47.07	40.54	45.28
Mmlu-pro	23.63	21.37	26.36	24.61	25.99	23.34	25.77	27.08
GPQA	30.30	8.08	28.79	36.87	28.79	9.60	27.78	28.79
Humaneval	41.46	40.85	34.76	37.20	37.20	41.46	35.98	39.63
Mbpp	23.80	18.00	0.40	6.00	20.80	26.40	3.60	27.80
Gsm8k	61.18	59.21	58.38	62.02	61.26	58.07	60.27	42.91
Math-500	49.06	43.16	46.96	43.68	43.00	44.84	48.48	44.62
TheoremQA	16.62	16.32	19.12	13.00	15.38	17.50	14.37	15.88
Cmmlu	45.36	45.99	45.00	45.88	46.01	46.44	44.84	46.60
Ceval	41.14	42.66	41.22	42.88	42.18	41.76	43.08	41.38

Table 7: Detailed staged experimental results on Qwen3-0.6B

Benchmark	4B	SFT	GRPO	S→G	S→A-G	A-SFT	A-GRPO	A-S→A-G
Avg	62.52	60.80	68.11	65.75	66.10	60.67	68.39	67.93
IFEVAL	78.19	77.82	87.06	88.91	90.02	81.15	88.54	90.02
MultiIF	82.00	83.37	86.74	90.58	90.88	83.79	87.80	91.28
Mmlu	69.90	72.78	70.02	71.47	70.96	71.91	70.00	69.86
Mmlu-pro	35.56	40.51	54.88	53.43	54.59	31.29	51.46	56.35
GPQA	30.81	10.10	35.35	22.22	19.19	13.13	41.41	27.27
Humaneval	75.61	76.83	78.66	79.88	82.32	79.27	81.10	81.71
Mbpp	65.00	60.40	64.20	63.20	61.80	61.40	65.20	64.80
Gsm8k	83.78	86.81	89.92	89.08	89.31	86.73	87.04	90.14
Math-500	70.26	62.20	70.14	63.36	62.66	63.72	70.7	65.92
TheoremQA	24.38	21.75	39.50	25.75	29.62	19.50	37.62	34.88
Cmmlu	72.09	72.18	72.03	72.48	72.63	71.80	72.07	72.45
Ceval	62.62	64.84	68.80	68.58	69.22	64.34	67.70	70.45

Table 8: Detailed staged experimental results on Qwen3-4B

Benchmark	0.6B Model			4B Model		
	SFT	CHORD	ACHORD	SFT	CHORD	ACHORD
Avg	40.81	41.19	42.09	62.52	69.30	69.27
IFEVAL	55.64	74.49	76.16	78.19	87.06	87.80
MultiIF	57.60	72.54	74.34	82.00	86.87	88.74
Mmlu	43.91	41.49	44.10	69.90	70.29	69.94
Mmlu-pro	23.63	25.53	25.33	35.56	55.90	56.04
GPQA	30.30	22.73	28.79	30.81	38.38	40.40
Humaneval	41.46	35.98	35.98	75.61	81.71	77.44
Mbpp	23.80	16.20	16.20	65.00	65.80	63.20
Gsm8k	61.18	55.57	55.57	83.78	87.41	90.86
Math-500	49.06	43.86	43.86	70.26	68.34	72.30
TheoremQA	16.62	16.12	16.12	24.38	47.25	41.25
Cmmlu	45.36	45.85	45.85	72.09	71.96	72.69
Ceval	41.14	43.94	43.94	62.62	70.63	70.62

Table 9: Detailed results of the unified experiments

Benchmark	SFT	0.1	0.01	0.001	0.0001	0.00001
Avg	43.97	38.18	47.53	47.83	47.67	47.61
IFEVAL	62.66	76.52	78.74	80.78	81.33	79.30
MultiIF	57.59	72.70	76.13	77.14	76.49	76.36
Mmlu	52.90	50.92	51.99	52.02	51.91	52.56
Mmlu-pro	27.88	24.89	28.19	27.38	27.42	26.72
GPQA	19.19	25.76	30.30	29.29	25.25	28.28
Humaneval	32.93	3.66	35.37	32.93	34.15	33.54
Mbpp	38.80	35.80	38.20	38.20	38.60	37.00
Gsm8k	62.70	26.99	62.77	63.38	65.81	64.75
Math-500	45.10	16.44	42.22	43.46	44.00	44.54
TheoremQA	23.00	21.61	22.50	24.88	22.75	23.75
Cmmlu	51.02	49.94	50.13	50.17	50.15	50.15
Ceval	53.84	52.94	53.87	54.34	54.21	54.21

Table 10: Detailed experimental results on GT coef

Benchmark	SFT	0.1	0.01	0.001	0.0001	0.00001
Avg	42.47	38.09	46.21	45.86	46.22	46.89
IFEVAL	57.12	75.97	78.93	78.74	78.74	78.19
MultiIF	59.53	73.09	76.49	75.70	75.42	75.62
Mmlu	52.82	49.81	52.82	53.09	53.17	53.47
Mmlu-pro	27.04	23.23	26.81	27.08	26.85	28.12
GPQA	10.10	28.28	30.30	19.70	27.27	24.75
Humaneval	40.85	22.56	34.76	38.41	34.15	38.40
Mbpp	35.00	28.60	36.60	36.40	35.80	34.80
Gsm8k	61.64	21.23	62.47	63.53	63.31	64.59
Math-500	38.16	10.68	29.16	29.86	34.58	36.74
TheoremQA	21.75	21.38	22.88	23.12	21.38	22.88
Cmmlu	50.91	49.93	50.66	50.51	50.81	50.83
Ceval	54.78	52.36	52.61	54.10	53.13	54.24

Table 11: Detailed experimental results on Teacher coef

Benchmark	RLVR	POLAR+RLVR	POLAR+A_{coef}	w A_{coef}	NoKL	REF	REFNoKL
Avg	47.52	46.70	47.68	47.83	47.52	46.25	40.65
IFEVAL	76.52	69.87	74.86	80.78	83.55	79.11	82.26
MultiIF	72.70	67.06	71.32	77.14	77.74	75.74	78.50
Mmlu	52.60	52.78	53.14	52.02	52.91	53.26	51.84
Mmlu-pro	27.89	29.00	29.48	27.38	27.90	27.02	26.71
GPQA	28.79	26.26	25.25	29.29	27.27	17.17	31.31
Humaneval	35.37	37.20	42.07	32.93	36.00	37.20	38.40
Mbpp	37.20	36.40	35.80	38.20	34.80	35.60	37.60
Gsm8k	64.29	66.72	66.03	63.38	63.91	65.05	7.20
Math-500	43.70	46.10	46.62	43.46	38.90	37.40	8.34
TheoremQA	23.50	25.37	22.62	24.88	22.88	23.50	22.38
Cmmlu	50.67	50.48	50.48	50.17	50.21	50.58	50.21
Ceval	53.77	53.16	54.54	54.34	53.03	53.32	53.03

Table 12: Detailed experimental results on other

B.4 Example of Training data

Eva Func Example

```python
class KeywordChecker(Instruction):
    def build_description(self, *, keywords=None):
        """Build the instruction description.
        Args:
            keywords: A sequence of strings representing the
                keywords that are expected in the response.
        Returns:
            A string representing the instruction description."""

        self.keywords = keywords
        self._description_pattern = "Include keywords {keywords} in
            the response."
        return self._description_pattern.format(keywords=keywords)

    def check_following(self, value):
        """Check if the response contain the expected keywords."""
        for keyword in self._keywords:
            if not re.search(keyword, value, flags=re.IGNORECASE):
                return False
        return True
```

Training Example

Question

Jamal works at a library shelving books. He has a cart full of books to put away in different sections. In the history section, he shelves 12 books. In the fiction section, he shelves 19 books. In the children's section, he shelves 8 books but finds 4 that were left in the wrong place that he adds to his cart to shelve elsewhere. He still has 16 books to shelve. How many books did he start with in the cart? Include a palindrome in your response. There should be 2 paragraphs. Paragraphs are separated by the markdown divider: *** Include keywords ['item', 'peak', 'population', 'team'] in the response. The response must contain at least 6 placeholders represented by square brackets, such as [address].

Ref Answer

Jamal started his day at the library, a place where the [population] of book lovers often gathers. His task was to shelve books that had been returned or misplaced. In the history section, he carefully placed 12 books on the shelves, ensuring each [item] was in its correct spot. Moving on to the fiction section, he shelved 19 more books, making sure that the novels were in alphabetical order. In the children's section, he shelved 8 books but found 4 that were misplaced, which he added back to his cart. After all this, he still had 16 books to shelve. To find out how many books Jamal started with, we need to add the books he shelved (12 + 19 + 8) to the books he found (4) and the books he still had to shelve (16). This gives us a total of [total_books] books.
*** Interestingly, as Jamal was shelving books, he noticed a unique book titled "Peak" by a lesser-known author. The book's title was a [palindrome], meaning it reads the same backward as forward. This intrigued him, and he made a mental note to read it later. As he continued his work, he realized that the library [team] was like a well-oiled machine, with everyone contributing to the smooth operation of the facility. By the end of his shift, Jamal felt a sense of accomplishment, knowing that he had played a crucial role in maintaining the library's organization. He started with [total_books] books in his cart, ensuring that each [item] found its rightful place on the shelves.

Teacher Answer

Jamal began his shift with a cart full of books, each book an item awaiting its proper place. He first visited the history section, where he shelved 12 books, followed by the fiction section, where he placed 19 more. In the children's section, he managed to shelve 8 books but discovered 4 misplaced ones, which he added back to his cart for later organization. Despite these comings and goings, he still has 16 books left to shelve. To determine how many books he started with, we calculate: 12 (history) + 19 (fiction) + 8 (children's) - 4 (added back) + 16 (remaining) = 51 books total. This number represents the peak effort required by Jamal's team to maintain order across the library's vast population of literary works. Each shelving decision impacts the accessibility of every item, much like how data at [address] influences urban population trends.
***The library's system relies on consistent workflows, similar to how a well-coordinated team manages resources at [facility_name]. The discovery of misplaced books highlights the importance of audits, especially during peak hours when the population of visitors surges. Jamal's cart, once full, now holds only 16 books—an improvement, but not yet complete. The palindrome "121" reflects the symmetry needed in such organizational tasks, where every action has a balanced reaction. Additional considerations include [city_name], [state_name], [zip_code], [phone_number], [email], and [department_head], all of which support the larger mission of knowledge distribution. Ultimately, each item shelved brings the library closer to optimal order, mirroring how accurate data at [address] supports demographic studies on population shifts.

Instruction ID
'keywords':'palindrome', 'paragraphs':'paragraphs',
'keywords':'existence', 'detectable_content':'number_placeholders'
Kwargs
None, {}, {'keywords': ['item', 'peak', 'population', 'team']}, {'num_placeholders': 6}

B.5 Example of answer between different training paradigm

Response Example

Instruction:
Please explain in detail the process of photosynthesis in plants, including its stages, key molecules, and overall importance. In your explanation, make sure to include the keywords ['chlorophyll', 'sunlight', 'glucose', 'oxygen'].

Teacher Answer (Good, Long, Correct)
Photosynthesis is a fundamental biological process by which green plants, algae, and certain bacteria convert light energy from the sun into chemical energy stored in glucose molecules. This process occurs primarily in the chloroplasts of plant cells, which contain the pigment chlorophyll that captures sunlight efficiently. Photosynthesis can be divided into two major stages: the light-dependent reactions and the Calvin cycle (light-independent reactions). During the light-dependent reactions, chlorophyll absorbs sunlight and drives the splitting of water molecules, releasing oxygen as a byproduct and generating energy carriers ATP and NADPH. In the Calvin cycle, carbon dioxide from the atmosphere is fixed into glucose using the ATP and NADPH generated in the first stage. This glucose serves as a primary energy source for the plant, while the oxygen released supports the respiration of other living organisms. Understanding this process is crucial for fields such as agriculture, ecology, and environmental science, as it underpins plant growth, food production, and global oxygen balance.

Model Answer (Incorrect, Short)
Photosynthesis is how plants get energy from sunlight.

SFT Model Answer (Incorrect, Long)
Photosynthesis occurs when plants absorb sunlight using chlorophyll and turn it into energy that they can use immediately. The process involves creating sugar in the leaves and distributing it throughout the plant for growth. While it generates glucose, the role of oxygen release and carbon fixation is misunderstood, making the explanation inaccurate.

GRPO Answer (Correct, Short)
Plants convert sunlight, water, and carbon dioxide into glucose and oxygen using chlorophyll.

Adversarial Answer (Correct, Medium-Long, Similar to Teacher)
Photosynthesis is the complex process through which green plants convert sunlight into chemical energy stored in glucose. Chlorophyll in the chloroplasts captures sunlight, which drives the splitting of water molecules, releasing oxygen as a byproduct and creating ATP and NADPH. During the Calvin cycle, carbon dioxide is fixed into glucose using these energy carriers. This process is vital not only for plant growth but also for maintaining atmospheric oxygen levels and supporting other life forms. The keywords 'chlorophyll', 'sunlight', 'glucose', and 'oxygen' are all included as requested, ensuring the explanation is complete and accurate.